

Course CSE 532: Multiprocessor Architecture

Authors: Rishabh Jain, Niramay Vaidya

Project Title: Performance Analysis of Association Rule Mining

Note: Both authors have contributed equally (equal division of work between the three implementations and corresponding analysis)

GitHub repository: https://github.com/rishucoding/532-perf_analysis


Note: CUDA code is available in branch `simple_bitmap` whereas the `pthread`s and serial code can be found in the `main` branch

Outline:

1. Quick pointers on the application
2. Dataset
 - 2.1. Structure
 - 2.2. Generation
 - 2.3. Parameters under control
 - 2.4. Practical relatability
3. Implementation specifics
 - 3.1. Naive sequential
 - 3.2. Parallelizing naive
 - 3.3. Smartly doing sequential
 - 3.4. Parallelizing smart
 - 3.5. Running on GPU
4. Correctness verification
5. Measurement study
 - 5.1. Timing measurement
 - 5.2. Varying `num_tx`
 - 5.3. Varying `num_items`
 - 5.4. Varying probability
 - 5.5. Varying threshold
 - 5.6. Special cases
 - 5.7. Summarizing the comparisons between different cases
 - 5.8. Summarizing the linux utilities used
6. Programming challenges
7. Hardware bottlenecks
8. Limitations and Future directions
9. References

1. Quick pointers on the application

[2] mentions that determining the buying patterns of demographic groups or segmenting customers according to relationships in their buying patterns is called data mining. A particular type of data mining is mining for associations. The goal here is to discover relationships / associations using the available information between different customers and their transactions and to generate rules for the inference of customer behaviour. In more concrete terms, this problem entails sets of 'k' items that are found to occur together in more than a certain threshold fraction of the transactions, given a database / dataset. Such a set of items would then be identified as a frequent itemset. In order to do this, a property that an itemset is frequent if the itemsets it was formed from (those in the previous level) are frequent can be utilized. Using this, the problem essentially boils down to performing computations in the manner shown below (taken from [1]), in order to exploit this property-



Grouping items using incremental lexicographic ordering at each level is called equivalence class partitioning. Such a partitioning helps ease the computation for determining the candidate itemsets at the next level.

Note: The only difference in our implementation is that instead of using alphabets to represent items, increasing positive integers have been used.

2. Dataset

2.1. Structure

A typedef struct is used to define an entry in the transaction database. This entry includes two fields: a txid to store the id of the entry, and an int array to mark all the items present. The upper limit on the total items is presently 100 and it is parameterized. The total number of entries is defined by a constant NUM_TX, and the total number of items is defined by a constant NUM_ITEMS. Note that the database can be visualized as a 2D array (matrix) where each row represents a transaction timestamp, and each column indicates the items purchased in that transaction.

2.2. Dataset generation

A method called “**generate_dataset**” is used which populates each cell in the matrix. A uniform probabilistic distribution is used to decide the chance of purchasing an item in any transaction entry. The value of chance is parameterized. Having a low chance gives a sparse matrix while a high value of chance gives a dense matrix.

For the special case datasets (having particular distributions representing skewness in data), a python script has been written to generate 4 different matrix variants of a particular size.

2.3. Parameters under control:

- a) NUM_TX - this represents the total entries in the database
- b) NUM_ITEMS - this represents the total columns/items in the database.
- c) THRESHOLD - this represents the minimum support
- d) Probability - For a given transaction entry, this determines the chance of purchasing an item.

2.4. Practical relatability (making sense in the real world)

1. During situations like off-season, holidays, high inflation, etc, the purchase of items could go down, and thus, each transaction would have a lower number of items marked. Thus, for this case, a sparse matrix is an appropriate representative.
2. On the other hand, when there is a high demand or right season for certain produce, like black friday deals, winter sales, etc, people often tend to buy in bulk, and stock up things. This would mean a high number of items marked per transaction. Thus, a dense matrix appropriately represents this case.
3. Threshold decides our association rule. For example, a 33% threshold means that we want to study all combinations of items which occur on average in every 3rd transaction. This parameter could help the sales expert in deciding what items are being frequently purchased, and thus give an opportunity for introducing varieties. Like, if milk is purchased frequently, a wise choice would be to introduce flavoured milk, almond milk, etc. However, if some item is not being purchased frequently, like steel utensils or study lamps, it could mean that these items serve for a long time.

3. Implementation specifics

3.1. Naive sequential (1 core)

As explained in [2], the naive way to solve this problem is to traverse the entire dataset to determine L1, by counting frequencies of all itemsets of size 1. From L1, the candidate itemset C2 can be computed. Using C2, the entire dataset can again be traversed to determine L2. This process can be continued as is further to generate all the frequent itemsets at each subsequent level. This algorithm has high computational complexity as it involves traversing the entire dataset for each candidate itemset to determine if it's frequent or not.

3.2. Parallelizing naive (multi-core)

According to [2], the way to parallelize this problem is to compute up to L2 serially, and then based on the equivalence classes present at the 2nd level, launch one thread for each equivalence class and beyond this point, each thread will keep on asynchronously generating frequent itemsets at each level. Every thread will perform the same mechanism henceforth as explained above in the naive sequential section.

3.3. Smartly doing sequential (1 core)

An algorithmic improvement over the naive sequential method is to compute as before up to L2, and then restructure the dataset such that it is now represented in the forms of lists of transactions for each itemset within L2. This saves traversing the entire original dataset for each subsequent level. This transformation of the dataset helps in combining the generation of the next level's candidate and further, frequent itemsets, from each previous level's frequent itemsets.

Steps:

1. A vector of vectors of int is used to represent each L_i (frequent itemsets at each level). The inner vector contains items in each itemset and the outer vector represents the total frequent itemsets at each level. For each level, the length of the inner vector increases by 1.
2. A vector of sets of int is used to represent the restructured dataset for each level. The inner sets contain the txids for each corresponding frequent itemset at that level.
3. A vector of indexes is required to be computed at the start of each level to determine the equivalence class boundaries in order to determine candidate itemsets from each equivalence class for the generation of the next level of frequent itemsets using this equivalence class.
4. To determine if a particular itemset at the next level is frequent, a set intersection operation is done on the txid sets of the corresponding current level's frequent itemsets using which this next level's itemset has been formed.
5. To form this next level's itemset in question, a set union operation needs to be performed.

6. Finally, the “**compute_li**” function iteratively performs the generation of the next level frequent itemsets and outputs the above mentioned data structures for the next level, by taking as input the previous level’s similar data structures.
7. C++’s STL (Standard Template Library) has been used extensively to perform all the operations mentioned above. The chrono utility has been used to measure the execution time.

Phases of development:

1. Initially, We figured out the computation of L2 and reformation of the dataset. (**commit ID- 7db323f**)
2. Following this, we determined what data structures would be needed which could generically be used for all levels and implemented the L3 computation. (**commit ID- f4d4a9a**)
3. Then we extended the logic for L3 to compute all frequent itemsets at all levels. (**commit ID- f058864**)
4. Finally, we added the measurement of execution time. (**commit ID- 28f0350**)
5. Once the smart algorithm was in place, we added support for the naive algorithm. (**commit ID- 400280c**)

3.4. Parallelizing smart (multi core)

The only difference with respect to the naive thread implementation here is that as per the smart sequential mechanism explained above, each thread has all the data now local to it (i.e. the required txids for the frequent itemsets corresponding to the equivalence class it is responsible for).

Phases of development:

1. Figuring out how to parallelize from L2 based on each generated equivalence class by assigning a thread to each was straightforward. Certain modifications involving the need of a target function to be provided to the pthread_create call were required, but most of the other logic remains the same as that present in the sequential case. (**commit ID- ae670fd**)
2. We then added support for setting cpu / core affinity to each thread in a round robin fashion based on the number of available cores. (**commit ID- 9372d67**)

3.5. Running on the GPU (accelerator)

The levels L1 and L2 are generated on the CPU. Further levels (L3, L4, ..., etc) are generated on the GPU. As mentioned above, in the CPU, the data for the levels are stored as vector of vector, and corresponding transaction ids are stored as a vector of sets. Since the input and output data sent to the GPU is a simple 1D array, there is an additional data processing step before and after the kernel launch. This data processing involves two parts: (1) linearizing the present levels vector’s to prepare the input data to the GPU kernel (2) converting the output array from the GPU into next levels vector’s.

Steps:

1. Each level is computed in one iteration of the for loop.
2. For each level, an index array is calculated by the CPU to identify all the equivalence classes.
3. GPU kernels are launched for each equivalence class in the present level, where a GPU thread runs a cuda kernel to evaluate one candidate pair. This allows all candidates to be evaluated in parallel.
4. If the candidate pair is valid, it's corresponding next level data is stored in the output array. This computation involves intersection and union operations on two sub arrays.

Phases of development:

1. After figuring out that the cuda kernel can compute each level one by one, the first choice was to launch only one kernel per level which works on one equivalence class. This allows all the equivalence classes to be evaluated in parallel. However, when larger database configurations were set, this approach requires a lot of data transfer between the devices which lead to segmentation faults. **(commit ID- f1d691bc)**
2. Further improvement in (1) was to change the indexing strategy when writing to output arrays. This helped in covering some large database configurations, but it still suffered from segmentation faults. **(commit ID- 101d6db)**
3. This motivated us to think of a new way to parallelize the computation. So, instead of parallelizing for all equivalence classes, it was decided to parallelize for all the candidates within one equivalence class. This also gives more parallel opportunities as many times one equivalence class has 1000's of candidates to evaluate. **(commit ID- b36c171)**
4. Features of unified virtual memory and asynchronous prefetch were used. **(commit ID- 2b224c1)**
5. **nvprof** profiling tool was used to study different metrics like data transfer size, cuda API calls with individual time taken, and number of CPU page faults.

4. Correctness verification

As a first step, an initial correctness verification of all the 3 implementations (sequential, pthreads and cuda) was done on a very small database (10 transactions, 8 items, threshold = 3). This was achieved by comparing the generated outputs of each level with a manually calculated version of the outputs.

Once this was successful, to verify correctness of the implementation for any given dataset, a comparison among output files was made. This includes file size and the total count of the generated frequent itemset, from which it was concluded that all implementations are in sync, since these values turned out to be the exact same.

5. Measurement study

5.1. Timing measurement


1. All the obtained execution time results have been averaged over 3 measurements to ensure more accuracy.
2. A stopwatch style start and stop measurement is taken using the chrono c++ library by surrounding the start/stop around any computation block-of-interest.

5.2. Varying NUM_TX (number of transactions / txids)

NUM_TX	NUM_FREQ_ITEMSETS	MAX_LEVEL_REACHED
50	62	L3
100	1680	L4
150	8240	L5
200	21317	L5
250	57913	L6
300	97825	L6
350	177942	L7
400	272667	L7
450	445848	L7
500	532578	L7

The above table shows the range of NUM_TX [50, 500] at a step size of 50 with the corresponding count of total frequent items sets and maximum level reached. The other parameters are fixed, particularly NUM_ITEMS = 50, probability = 50%, and threshold = 20. It can be seen that with an increase in the number of transactions, the total frequent sets and max level increases. This is because more transactions gives opportunity for more associations.

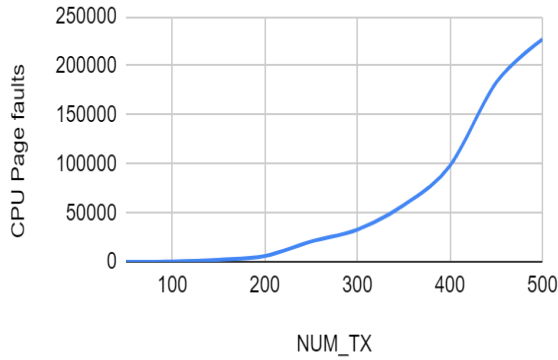
Execution Time vs NUM_TX




The above figure shows the variation in execution time as the number of transactions increases. It can be seen that the execution time increases with increase in NUM_TX. Both multi core and GPU perform much better than single core. At NUM_TX = 500, both multi core and GPU perform 4 times better than single core. There is not much performance difference between multi core and GPU. This could be because of the multi core implementation taking good advantage of 48 CPU cores in ladon server.

The next three figures discuss the profiling of CUDA implementation.

CPU Page faults vs. NUM_TX




Data Transfer vs NUM_TX



The above figure shows the variation in CPU Page faults and Data transfer as the number of transactions increases. Since unified memory is used, CPU Page faults occur whenever GPU accesses memory. The device to host memory size is higher than host to device because the output array size considers all candidates to be counted for the next level. Also, many times, the outputs are spatially far, thus falling into different pages, and leading to more data transfer. It is observed that the data transfer size is very high (like in GBs) when NUM_TX is large.

Data processing time vs. NUM_TX



The above figure shows the variation in data processing time as the number of transactions increases. This processing is the same as mentioned in section 3.5 which happens before and after every kernel launch. In the present implementation, this is actually an overhead to use

GPU. We believe this can be reduced significantly by cleverly designing the input and output parameters for the CUDA kernel.


5.3. Varying NUM_ITEMS (number of items)

Number of frequent itemsets and the highest level reached with varying number of items (along with the thread count in case of the pthreads implementation, based upon number of equivalence classes at level 2)-

NUM_ITEMS	NUM_FREQ_ITEMSETS	MAX_LEVEL_REACHED	THREAD_COUNT
10	47	L3	9
20	211	L4	19
30	532	L4	29
40	1032	L4	38
50	1400	L4	49
60	2591	L5	59
70	2664	L4	69
80	5206	L5	79
90	5704	L4	89
100	9521	L5	99

The above table shows the range of NUM_ITEMS [10, 100] at a step size of 10 with the corresponding count of total frequent items sets, maximum level reached and thread count. The other parameters are fixed, particularly NUM_TX = 100, probability = 50%, and threshold = 20. It can be seen that with an increase in the number of items, the total frequent sets, max level reached and thread count increases. Similar to the previous table, this is because more items per transaction gives better opportunities for more associations. Also, with an increase in the number of items, there is a proportionate increase in the number of threads launched. This interesting observation is not seen in other cases / variations.

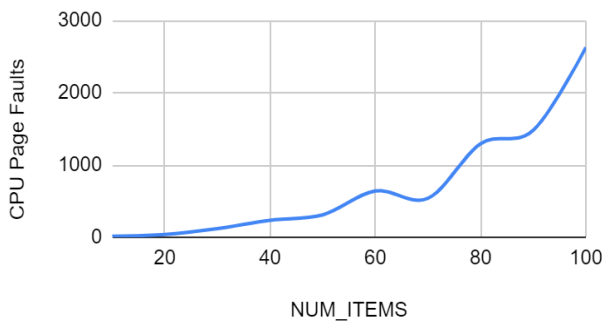
Execution Time vs NUM_ITEMS




The above figure shows the variation in execution time as the number of items increases. It can be seen that the execution time increases with increase in NUM_ITEMS. Both multi core and GPU perform much better than single core. At NUM_TX = 100, GPU performs 10.8 times better over single core, while multicore performs 6.6 times better over single core. As NUM_ITEMS increase, a notable performance difference is observed between multicore and GPU. This is because more items generates more candidates which our CUDA kernel evaluates.

The next three figures discuss the profiling of CUDA implementation.

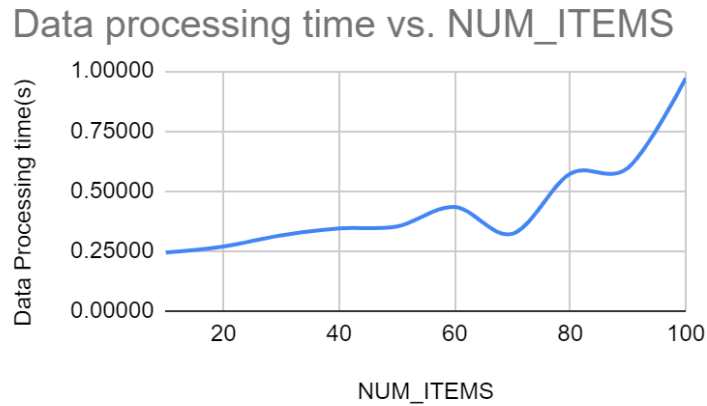
CPU Page Faults vs. NUM_ITEMS



Data Transfer vs NUM_ITEMS



The above figure shows the variation in CPU Page faults and Data transfer as the number of items increases. The data transfer size doesn't reach GBs unlike previous cases. This is because the number of items are incremented by step size of 10 and not 50.



The above figure shows the variation in data processing time as the number of items increases. In the above three figures, the strange dips at 70 and 90 could be because of the distribution of the sample chosen. It could be corrected when the performed analysis is averaged over various samples of the same size.

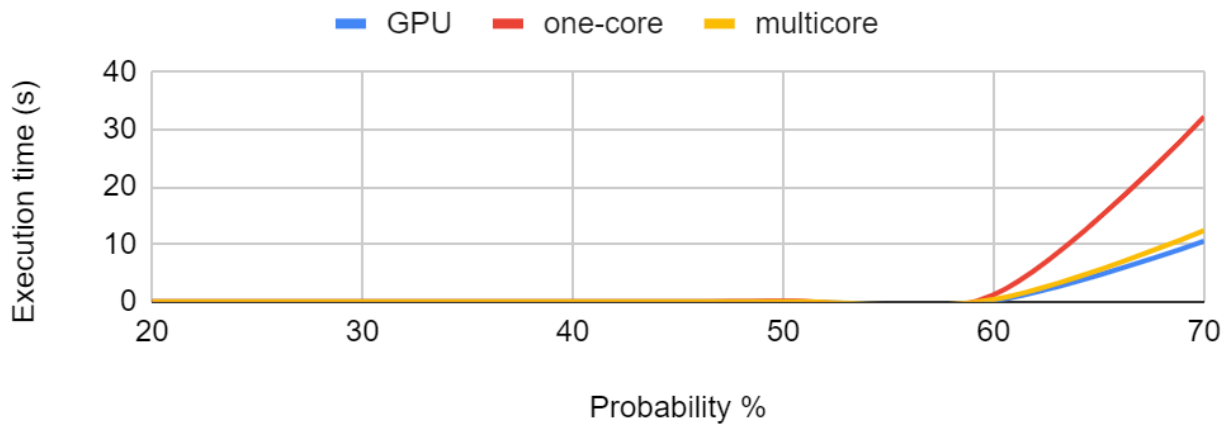
5.4. Varying probability

Number of frequent itemsets and the highest level reached with varying probability-

PROBABILITY	NUM_FREQ_ITEMSETS	MAX_LEVEL_REACHED
20%	23	L2
30%	49	L2
40%	312	L3
50%	1454	L5
60%	18251	L6
70%	699349	L9

The above table shows the range of probability [0.2, 0.7] at a step size of 0.1 with the corresponding count of total frequent items sets and maximum level reached. The other parameters are fixed, particularly NUM_TX = 100, NUM_ITEMS = 50, and threshold = 20. It can be seen that with an increase in the probability, the total frequent sets and max level reached increases. Similar to the previous table, this is because a higher probability per transaction gives better opportunity for more associations.


Execution Time vs Probability %




The above figure shows the variation in execution time as the % probability. It can be seen that the execution time exponentially increases. Both multi core and GPU perform much better than single core for a probability > 40%. At probability = 70%, GPU performs 3 times better over single core, while multicore performs 2.6 times better over single core. As the probability increases, the execution time of all implementations exponentially increases. This is because the computation complexity (association opportunity) increases with a higher value of probability.

The next three figures discuss the profiling of CUDA implementation.

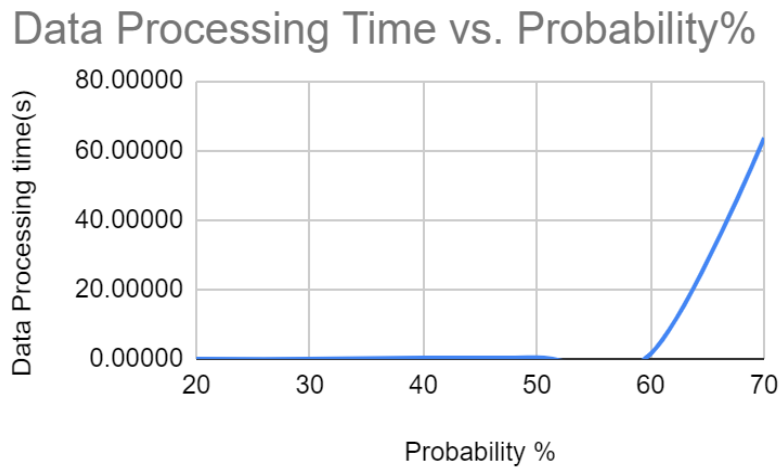
CPU Page Faults vs. Probability %



Data Transfer vs Probability



The above figure shows the variation in CPU Page faults and Data transfer as the probability increases.



The above figure shows the variation in data processing time as the probability increases. In all the three figures, the value in the Y axis tends to exponentially increase after 60% probability.


5.5. Varying threshold

Number of frequent itemsets and the highest level reached with varying threshold-

THRESHOLD	NUM_FREQ_ITEMSETS	MAX_LEVEL_REACHED
25	245898	L7
50	19749	L5
75	2429	L4
100	1251	L4
125	660	L3
150	61	L3
175	50	L2
200	50	L2

The above table shows the range of THRESHOLD [25, 200] at a step size of 25 with the corresponding count of total frequent items sets, maximum level reached and thread count. The other parameters are fixed, particularly NUM_TX = 500, NUM_ITEMS = 50, and probability = 50%. The values of threshold relates to % occurrence in NUM_TX, which is [5% to 40%]. It can be seen that with an increase in the threshold, the total frequent item sets and max level reached decreases. This is because a higher value of threshold is analogous to a strong filter, which decreases the scope of association.


Execution Time vs %Th



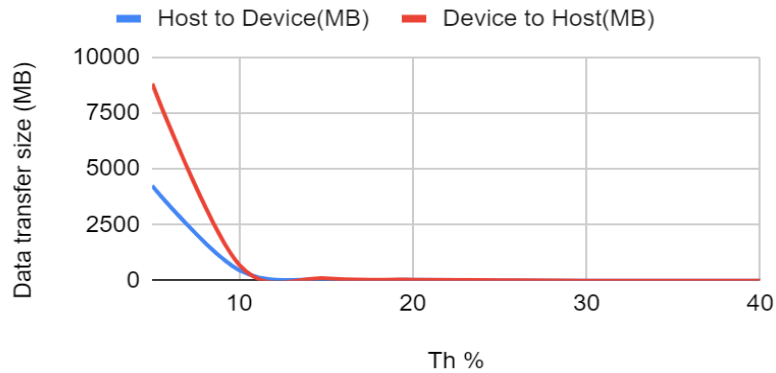
The above figure shows the variation in execution time as the % threshold increases. It can be seen that the execution time decreases with increase in threshold. Both multi core and GPU perform much better than single core for a lower value of threshold. At Threshold = 5%, GPU performs 6.5 times better over single core, while multicore performs 3.7 times better over single core. As the threshold increases, the execution time of all implementations converges. This is because the computation complexity (association opportunity) decreases with a higher value of threshold.

The next three figures discuss the profiling of CUDA implementation.

CPU Page Faults vs. Th%




Data Transfer vs Th %



The above figures show the variation in CPU Page faults and Data transfer as the threshold increases.

Data Processing Time vs. Th%



The above figure shows the variation in data processing time as the Threshold increases. In all the three figures, the value in the Y axis tends to exponentially decrease after the 12% threshold.

5.6. Special cases

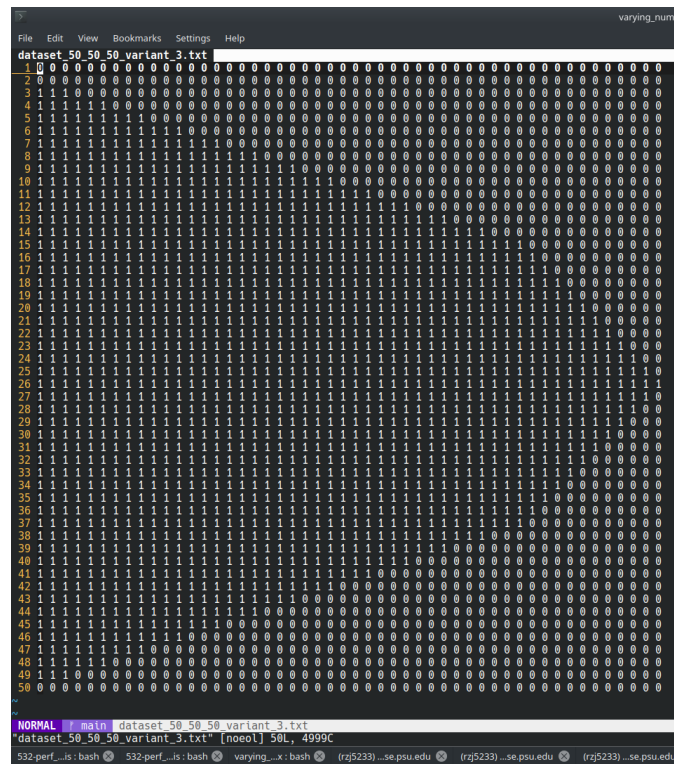
In order to study how the skewness in data affects the results in comparison to a random distribution, 4 different variants with the following configuration were generated-

NUM_TX = 50

NUM_ITEMS = 50


The baseline randomly distributed dataset is as follows-

This variant is analogous to a gaussian / normal distribution. Item popularity again decreases from left to right, the leftmost item being the most popular.



This variant is analogous to an exponential distribution. Item popularity again decreases from left to right, the rightmost item being the least popular.

Execution time pthreads vs NUM_TX



As seen in the graph above, there is almost no difference in terms of the naive implementation when core affinity is set or unset. In terms of the smart implementation, setting the cpu affinity helps since the execution time reduces for higher transaction sizes as compared to when the cpu affinity is unset. Also, as can be observed here, there comes a crossover point after which the smart implementation performs better than the naive implementation. Before this crossover point, the cost of maintaining the extra txids data structure exceeds that of exhaustively searching through the entire dataset.

The sar command line utility was used to record the cpu usage characteristics (it is similar to top and htop and reports the same statistics).

```
sar -P ALL <interval> <num_of_times> > <output_file>
```

Example: `sar -P ALL 1 500 > <output_file>`

CPU usage pattern with cpu affinity unset vs set-

FileEditViewBookmarksSettingsHelp

2 cpu_usage_affinity_set.txt

8672

8673 01:18:45 PM

CPU

Nuser

hnic

hssystem

hswaitat

hstotal

hslide

8674 01:18:46 PM

all

12.94

0.00

0.07

0.00

0.00

89.19

8675 01:18:46 PM

0

12.37

0.15

0.00

0.00

82.47

0.00

8676 01:18:46 PM

1

21.05

0.00

7.37

0.00

0.00

71.58

8677 01:18:46 PM

2

0.00

0.00

0.00

0.00

100.00

0.00

8678 01:18:46 PM

3

5.45

0.00

1.01

0.00

0.00

93.54

8679 01:18:46 PM

4

9.18

0.00

6.12

0.00

0.00

84.69

8680 01:18:46 PM

5

0.00

0.00

0.00

0.00

100.00

0.00

8681 01:18:46 PM

6

0.00

0.00

0.00

0.00

100.00

0.00

8682 01:18:46 PM

7

0.00

0.00

0.00

0.00

100.00

0.00

8683 01:18:46 PM

8

3.09

0.00

2.06

0.00

0.00

94.85

8684 01:18:46 PM

9

0.00

0.00

0.00

0.00

100.00

0.00

8685 01:18:46 PM

10

0.00

0.00

0.00

0.00

100.00

0.00

8686 01:18:46 PM

11

32.99

0.00

7.22

0.00

0.00

59.79

8687 01:18:46 PM

12

9.57

0.00

6.38

0.00

0.00

84.04

8688 01:18:46 PM

13

14.50

0.00

10.42

0.00

0.00

72.08

8689 01:18:46 PM

14

14.09

0.00

11.70

0.00

0.00

73.48

8690 01:18:46 PM

15

0.00

0.00

0.00

0.00

100.00

0.00

8691 01:18:46 PM

16

12.37

0.00

6.39

0.00

0.00

81.44

8692 01:18:46 PM

17

11.83

0.00

9.68

0.00

0.00

76.49

8693 01:18:46 PM

18

0.00

0.00

0.00

0.00

100.00

0.00

8694 01:18:46 PM

19

17.53

0.00

9.28

0.00

0.00

72.26

8695 01:18:46 PM

20

45.36

0.00

6.19

0.00

0.00

48.45

8696 01:18:46 PM

21

27.37

0.00

7.37

0.00

0.00

65.26

8697 01:18:46 PM

22

19.35

0.00

10.75

0.00

0.00

69.09

8698 01:18:46 PM

23

17.09

0.00

10.93

0.00

0.00

71.58

8699 01:18:46 PM

24

15.00

0.00

1.00

0.00

84.00

8700 01:18:46 PM

25

15.79

0.00

8.42

0.00

0.00

75.79

8701 01:18:46 PM

26

6.00

0.00

0.00

0.00

100.00

0.00

8702 01:18:46 PM

27

0.00

0.00

0.00

0.00

100.00

0.00

8703 01:18:46 PM

28

0.00

0.00

0.00

0.00

100.00

0.00

8704 01:18:46 PM

29

21.96

0.00

7.04

0.00

0.00

70.84

8705 01:18:46 PM

30

35.79

0.00

6.32

0.00

0.00

57.89

8706 01:18:46 PM

31

36.79

0.00

63.64

0.00

0.00

0.00

8707 01:18:46 PM

32

11.22

0.00

7.14

0.00

0.00

61.63

8708 01:18:46 PM

33

17.20

0.00

8.68

0.00

0.00

74.19

8709 01:18:46 PM

34

0.00

0.00

0.00

0.00

100.00

0.00

8710 01:18:46 PM

35

15.79

0.00

6.32

0.00

0.00

77.89

8711 01:18:46 PM

36

37.00

0.00

63.00

0.00

0.00

0.00

8712 01:18:46 PM

37

20.21

0.00

7.45

0.00

0.00

72.34

8713 01:18:46 PM

38

19.59

0.00

9.28

0.00

0.00

71.15

8714 01:18:46 PM

39

2.00

0.00

1.00

0.00

97.00

0.00

8715 01:18:46 PM

40

16.16

0.00

7.87

0.00

0.00

76.77

8716 01:18:46 PM

41

9.00

0.00

3.00

0.00

88.00

0.00

8717 01:18:46 PM

42

18.95

0.00

6.42

0.00

0.00

72.63

8718 01:18:46 PM

43

27.27

0.00

3.83

0.00

0.00

69.70

8719 01:18:46 PM

44

0.00

0.00

0.00

0.00

100.00

0.00

8720 01:18:46 PM

45

5.21

0.00

6.25

0.00

0.00

80.56

8721 01:18:46 PM

46

0.00

0.00

0.00

0.00

100.00

0.00

8722 01:18:46 PM

47

26.04

0.00

3.12

0.00

0.00

76.83

8723

8724 01:18:46 PM

CPU

Nuser

hnic

hssystem

hswaitat

hstotal

hslide

8725 01:18:47 PM

all

15.54

0.00

6.84

0.02

0.00

78.39

8726 01:18:47 PM

0

30.77

0.00

10.99

0.00

0.00

59.24

8673 01:18:45 PM8674 01:18:46 PM8675 01:18:46 PM8676 01:18:46 PM8677 01:18:46 PM8678 01:18:46 PM8679 01:18:46 PM8680 01:18:46 PM8681 01:18:46 PM8682 01:18:46 PM8683 01:18:46 PM8684 01:18:46 PM8685 01:18:46 PM8686 01:18:46 PM8687 01:18:46 PM8688 01:18:46 PM8689 01:18:46 PM8690 01:18:46 PM8691 01:18:46 PM8692 01:18:46 PM8693 01:18:46 PM8694 01:18:46 PM8695 01:18:46 PM8696 01:18:46 PM8697 01:18:46 PM8698 01:18:46 PM8699 01:18:46 PM8700 01:18:46 PM8701 01:18:46 PM8702 01:18:46 PM8703 01:18:46 PM8704 01:18:46 PM8705 01:18:46 PM8706 01:18:46 PM8707 01:18:46 PM8708 01:18:46 PM8709 01:18:46 PM8710 01:18:46 PM8711 01:18:46 PM8712 01:18:46 PM8713 01:18:46 PM8714 01:18:46 PM8715 01:18:46 PM8716 01:18:46 PM8717 01:18:46 PM8718 01:18:46 PM8719 01:18:46 PM8720 01:18:46 PM8721 01:18:46 PM8722 01:18:46 PM87238724 01:18:46 PM8725 01:18:47 PM8726 01:18:47 PM87278728872987308731873287338734873587368737873887398740874187428743874487458746874787488749875087518752875387548755875687578758875987608761876287638764876587668767876887698770877187728773877487758776877787788779878087818782878387848785878687878788878987908791879287938794879587968797879887998800

8673 01:18:45 PM8674 01:18:46 PM8675 01:18:46 PM8676 01:18:46 PM8677 01:18:46 PM8678 01:18:46 PM8679 01:18:46 PM8680 01:18:46 PM8681 01:18:46 PM8682 01:18:46 PM8683 01:18:46 PM8684 01:18:46 PM8685 01:18:46 PM8686 01:18:46 PM8687 01:18:46 PM8688 01:18:46 PM8689 01:18:46 PM8690 01:18:46 PM8691 01:18:46 PM8692 01:18:46 PM8693 01:18:46 PM8694 01:18:46 PM8695 01:18:46 PM8696 01:18:46 PM8697 01:18:46 PM8698 01:18:46 PM8699 01:18:46 PM8700 01:18:46 PM8701 01:18:46 PM8702 01:18:46 PM8703 01:18:46 PM8704 01:18:46 PM8705 01:18:46 PM8706 01:18:46 PM8707 01:18:46 PM8708 01:18:46 PM8709 01:18:46 PM8710 01:18:46 PM8711 01:18:46 PM8712 01:18:46 PM8713 01:18:46 PM8714 01:18:46 PM8715 01:18:46 PM8716 01:18:46 PM8717 01:18:46 PM8718 01:18:46 PM8719 01:18:46 PM8720 01:18:46 PM8721 01:18:46 PM8722 01:18:46 PM87238724 01:18:46 PM8725 01:18:47 PM8726 01:18:47 PM8727872887298730873101:18:45 PM01:18:46 PM


```
affinity_results: vi — Console
File Edit View Bookmarks Settings Help
cpu_usage_affinity_unset.txt
2195
2196 01:16:38 PM CPU %user %nice %system %iowait %steal %idle
2197 01:16:39 PM all 3.67 0.00 2.67 0.00 0.00 93.67
2198 01:16:39 PM 0 0.00 0.00 0.00 0.00 0.00 100.00
2199 01:16:39 PM 1 0.00 0.00 0.00 0.00 0.00 100.00
2200 01:16:39 PM 2 0.00 0.00 0.00 0.00 0.00 100.00
2201 01:16:39 PM 3 0.00 0.00 0.00 0.00 0.00 100.00
2202 01:16:39 PM 4 0.00 0.00 0.00 0.00 0.00 100.00
2203 01:16:39 PM 5 0.00 0.00 0.00 0.00 0.00 100.00
2204 01:16:39 PM 6 100.00 0.00 0.00 0.00 0.00 0.00
2205 01:16:39 PM 7 0.00 0.00 0.00 0.00 0.00 100.00
2206 01:16:39 PM 8 0.00 0.00 0.00 0.00 0.00 100.00
2207 01:16:39 PM 9 0.00 0.00 0.00 0.00 0.00 100.00
2208 01:16:39 PM 10 0.00 0.00 0.00 0.00 0.00 100.00
2209 01:16:39 PM 11 0.00 0.00 0.00 0.00 0.00 100.00
2210 01:16:39 PM 12 1.00 0.00 2.00 0.00 0.00 97.00
2211 01:16:39 PM 13 0.00 0.00 0.00 0.00 0.00 100.00
2212 01:16:39 PM 14 0.00 0.00 0.00 0.00 0.00 100.00
2213 01:16:39 PM 15 0.00 0.00 0.00 0.00 0.00 100.00
2214 01:16:39 PM 16 0.00 0.00 0.00 0.00 0.00 100.00
2215 01:16:39 PM 17 0.00 0.00 0.00 0.00 0.00 100.00
2216 01:16:39 PM 18 0.00 0.00 0.00 0.00 0.00 100.00
2217 01:16:39 PM 19 0.00 0.00 0.00 0.00 0.00 100.00
2218 01:16:39 PM 20 0.00 0.00 0.00 0.00 0.00 100.00
2219 01:16:39 PM 21 0.00 0.00 0.00 0.00 0.00 100.00
2220 01:16:39 PM 22 0.00 0.00 0.00 0.00 0.00 100.00
2221 01:16:39 PM 23 0.00 0.00 0.00 0.00 0.00 100.00
2222 01:16:39 PM 24 0.00 0.00 0.00 0.00 0.00 100.00
2223 01:16:39 PM 25 0.00 0.00 0.00 0.00 0.00 100.00
2224 01:16:39 PM 26 0.00 0.00 0.00 0.00 0.00 100.00
2225 01:16:39 PM 27 0.00 0.00 0.00 0.00 0.00 100.00
2226 01:16:39 PM 28 0.00 0.00 0.00 0.00 0.00 100.00
2227 01:16:39 PM 29 0.00 0.00 0.00 0.00 0.00 100.00
2228 01:16:39 PM 30 0.00 0.00 0.00 0.00 0.00 100.00
2229 01:16:39 PM 31 38.00 0.00 62.00 0.00 0.00 0.00
2230 01:16:39 PM 32 0.00 0.00 0.00 0.00 0.00 100.00
2231 01:16:39 PM 33 0.00 0.00 0.00 0.00 0.00 100.00
2232 01:16:39 PM 34 0.00 0.00 0.00 0.00 0.00 100.00
2233 01:16:39 PM 35 0.00 0.00 0.00 0.00 0.00 100.00
2234 01:16:39 PM 36 36.00 0.00 64.00 0.00 0.00 0.00
2235 01:16:39 PM 37 0.00 0.00 0.00 0.00 0.00 100.00
2236 01:16:39 PM 38 0.00 0.00 0.00 0.00 0.00 100.00
2237 01:16:39 PM 39 0.00 0.00 0.00 0.00 0.00 100.00
2238 01:16:39 PM 40 0.00 0.00 0.00 0.00 0.00 100.00
2239 01:16:39 PM 41 1.00 0.00 0.00 0.00 0.00 99.00
2240 01:16:39 PM 42 0.00 0.00 0.00 0.00 0.00 100.00
2241 01:16:39 PM 43 0.00 0.00 0.00 0.00 0.00 100.00
2242 01:16:39 PM 44 0.00 0.00 0.00 0.00 0.00 100.00
2243 01:16:39 PM 45 0.00 0.00 0.00 0.00 0.00 100.00
2244 01:16:39 PM 46 0.00 0.00 0.00 0.00 0.00 100.00
2245 01:16:39 PM 47 0.00 0.00 0.00 0.00 0.00 100.00
2246
NORMAL | :main| cpu_usage_affinity_unset.txt text utf-8(unix) 112,007 words 17% = 2246/12702 ln : 1
532-perf_analysis: bash 532-perf_analysis: bash varying_num_tx: bash (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu affinity_results: vi
```

```
affinity_results: vi — Console
File Edit View Bookmarks Settings Help
cpu_usage_affinity_unset.txt
3419
3420 01:17:02 PM CPU %user %nice %system %iowait %steal %idle
3421 01:17:03 PM all 3.73 0.00 2.54 0.00 0.00 93.73
3422 01:17:03 PM 0 0.00 0.00 0.00 0.00 0.00 100.00
3423 01:17:03 PM 1 0.00 0.00 0.00 0.00 0.00 100.00
3424 01:17:03 PM 2 0.00 0.00 0.00 0.00 0.00 100.00
3425 01:17:03 PM 3 0.00 0.00 0.00 0.00 0.00 100.00
3426 01:17:03 PM 4 1.00 0.00 0.00 0.00 0.00 99.00
3427 01:17:03 PM 5 0.00 0.00 0.00 0.00 0.00 100.00
3428 01:17:03 PM 6 0.00 0.00 0.00 0.00 0.00 100.00
3429 01:17:03 PM 7 0.00 0.00 0.00 0.00 0.00 100.00
3430 01:17:03 PM 8 0.00 0.00 0.00 0.00 0.00 100.00
3431 01:17:03 PM 9 0.00 0.00 0.00 0.00 0.00 100.00
3432 01:17:03 PM 10 100.00 0.00 0.00 0.00 0.00 0.00
3433 01:17:03 PM 11 0.00 0.00 0.00 0.00 0.00 100.00
3434 01:17:03 PM 12 0.00 0.00 0.00 0.00 0.00 100.00
3435 01:17:03 PM 13 0.00 0.00 0.00 0.00 0.00 100.00
3436 01:17:03 PM 14 0.00 0.00 0.00 0.00 0.00 100.00
3437 01:17:03 PM 15 0.00 0.00 0.00 0.00 0.00 100.00
3438 01:17:03 PM 16 0.00 0.00 0.00 0.00 0.00 100.00
3439 01:17:03 PM 17 0.00 0.00 0.00 0.00 0.00 100.00
3440 01:17:03 PM 18 0.00 0.00 0.00 0.00 0.00 100.00
3441 01:17:03 PM 19 0.00 0.00 0.00 0.00 0.00 100.00
3442 01:17:03 PM 20 0.00 0.00 0.00 0.00 0.00 100.00
3443 01:17:03 PM 21 0.00 0.00 0.00 0.00 0.00 100.00
3444 01:17:03 PM 22 0.00 0.00 0.00 0.00 0.00 100.00
3445 01:17:03 PM 23 0.00 0.00 0.00 0.00 0.00 100.00
3446 01:17:03 PM 24 0.00 0.00 0.00 0.00 0.00 100.00
3447 01:17:03 PM 25 0.00 0.00 0.00 0.00 0.00 100.00
3448 01:17:03 PM 26 0.00 0.00 0.00 0.00 0.00 100.00
3449 01:17:03 PM 27 0.00 0.00 0.00 0.00 0.00 100.00
3450 01:17:03 PM 28 0.00 0.00 0.00 0.00 0.00 100.00
3451 01:17:03 PM 29 0.00 0.00 0.00 0.00 0.00 100.00
3452 01:17:03 PM 30 0.00 0.00 0.00 0.00 0.00 100.00
3453 01:17:03 PM 31 40.00 0.00 60.00 0.00 0.00 0.00
3454 01:17:03 PM 32 0.00 0.00 0.00 0.00 0.00 100.00
3455 01:17:03 PM 33 0.00 0.00 0.00 0.00 0.00 100.00
3456 01:17:03 PM 34 0.00 0.00 0.00 0.00 0.00 100.00
3457 01:17:03 PM 35 0.00 0.00 0.00 0.00 0.00 100.00
3458 01:17:03 PM 36 38.38 0.00 61.62 0.00 0.00 0.00
3459 01:17:03 PM 37 0.00 0.00 0.00 0.00 0.00 100.00
3460 01:17:03 PM 38 0.00 0.00 0.00 0.00 0.00 100.00
3461 01:17:03 PM 39 0.00 0.00 0.00 0.00 0.00 100.00
3462 01:17:03 PM 40 0.00 0.00 0.00 0.00 0.00 100.00
3463 01:17:03 PM 41 0.99 0.00 0.00 0.00 0.00 99.01
3464 01:17:03 PM 42 0.00 0.00 0.00 0.00 0.00 100.00
3465 01:17:03 PM 43 0.00 0.00 0.00 0.00 0.00 100.00
3466 01:17:03 PM 44 0.00 0.00 0.00 0.00 0.00 100.00
3467 01:17:03 PM 45 0.00 0.00 0.00 0.00 0.00 100.00
3468 01:17:03 PM 46 0.00 0.00 0.00 0.00 0.00 100.00
3469 01:17:03 PM 47 0.00 0.00 0.00 0.00 0.00 100.00
3470
NORMAL | :main| cpu_usage_affinity_unset.txt text utf-8(unix) 112,007 words 27% = 3470/12702 ln : 1
532-perf_analysis: bash 532-perf_analysis: bash varying_num_tx: bash (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu (rj5233) ladon.cse.psu.edu affinity_results: vi
```

Comparison of naive and smart serial implementations-


Varying number of transactions

Execution time serial vs NUM_TX




Varying number of items

Execution time serial vs NUM_ITEMS




Varying threshold

Execution time serial vs THRESHOLD



Varying probability

Execution time serial vs PROBABILITY




Comparison of naive and smart pthreads implementations-

Note: For a varying number of transactions, the graph has already been explained before.


Varying number of items

Execution time pthreads vs NUM_ITEMS



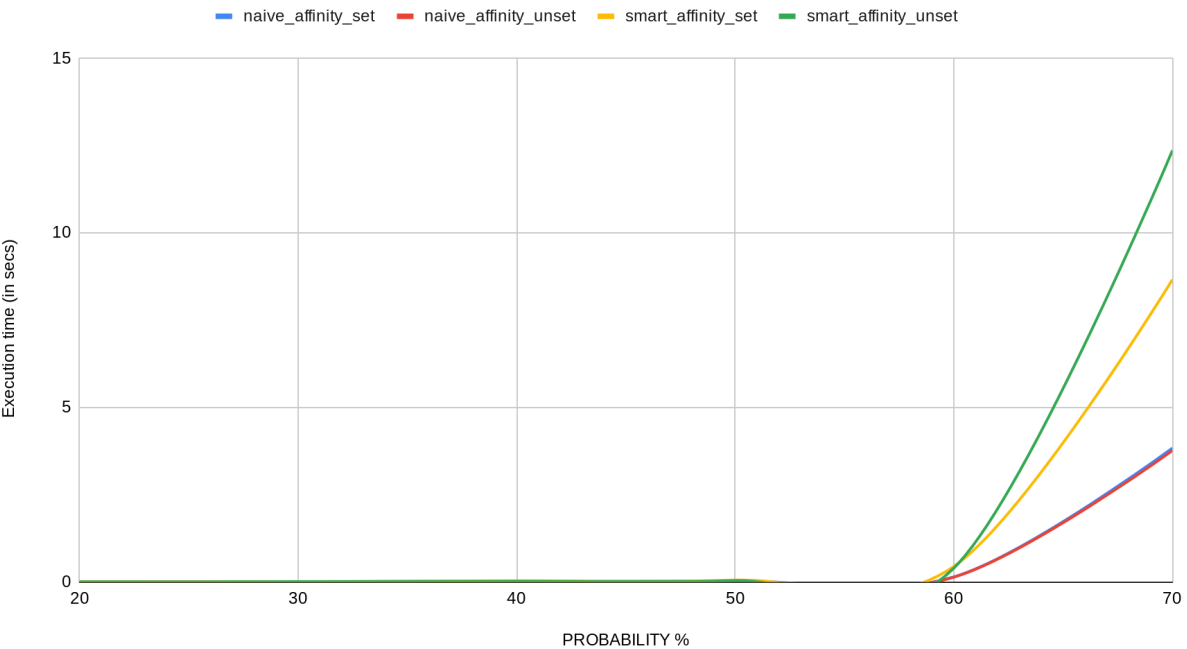
Varying threshold

Execution time pthreads vs THRESHOLD



Varying probability


Execution time pthreads vs PROBABILITY



Comparison of naive serial and pthreads implementations-


Varying number of transactions

Execution time naive vs NUM_TX



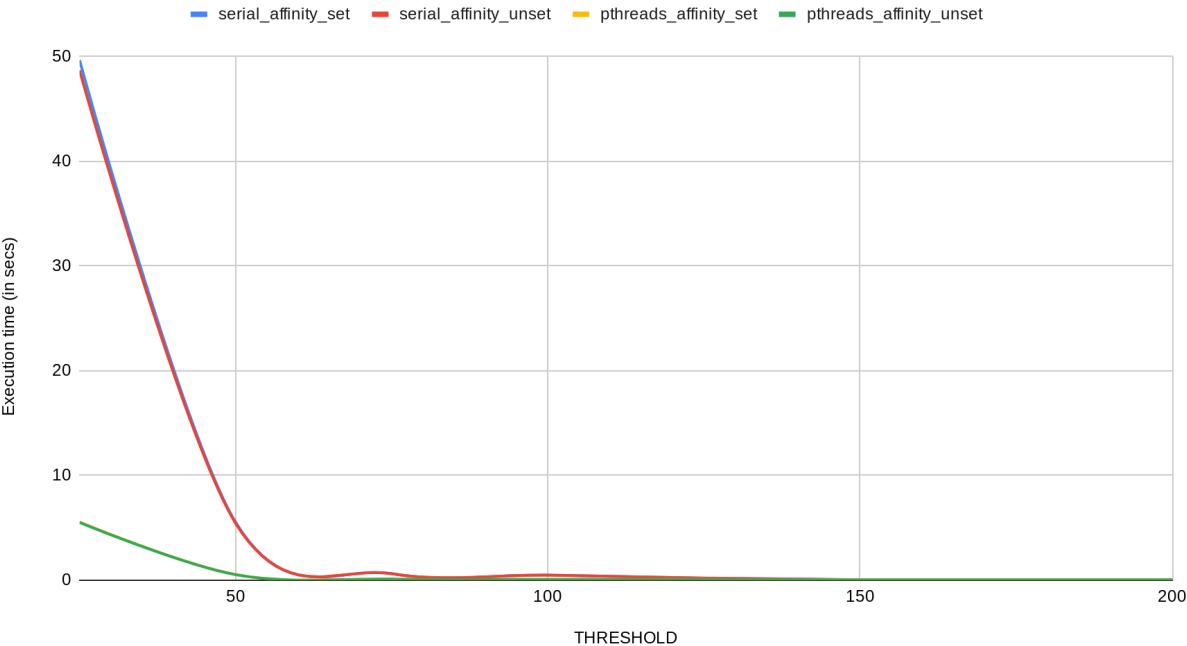
Varying number of items

Execution time naive vs NUM_ITEMS



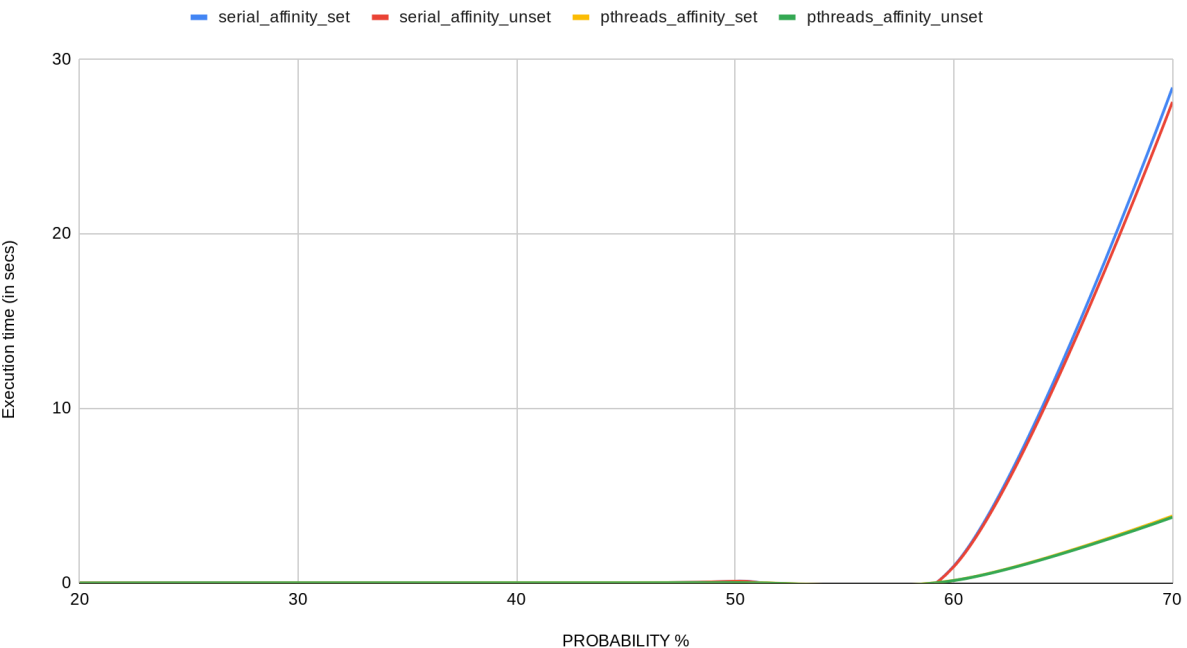
Varying threshold

Execution time naive vs THRESHOLD



Varying probability


Execution time naive vs PROBABILITY



Comparison of smart serial and pthreads implementations-


Varying number of transactions

Execution time smart vs NUM_TX




Varying number of items

Execution time smart vs NUM_ITEMS




Varying threshold

Execution time smart vs THRESHOLD



Varying probability

Execution time smart vs PROBABILITY



As observed in all the graphs above, the pthreads implementation performs significantly better than the serial implementation, irrespective of whether cpu affinity is set or unset. Further, there is no significant difference with respect to core affinity being set or not. Finally, in most of the above cases except in case of varying transactions for the serial implementation (smart performs better in this case), naive seems to perform better than smart (with a few minor anomalies in some cases, execution time vs threshold, for example). A possible reason for this has been explained in the report above (see the section of cpu affinity set and unset), though further analysis is required to get a complete understanding of this scenario.

Results obtained for higher number of transactions-

NUM_ITEMS = 50
 THRESHOLD = 20
 PROBABILITY = 0.5

num_tx	exec_time (in secs)							
	naive_serial		smart_serial		naive_pthreads		smart_pthreads	
	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset
500	82.89143 3	81.30213 3	32.85263 3	32.62836 7	9.539257	9.493583	7.94909	11.363633
750	500.224	496.788	128.441	129.754	55.8848	55.9105	62.2076	61.7823

1000	1286.22	1257.01	279.091	279.9	149.273	147.895	111.15	170.568
------	---------	---------	---------	-------	---------	---------	--------	---------

The following tabulated sequential measurement of time (the last 4 columns in tables below) is not an introduction of some new algorithm or method to perform the task at hand, but is just a measurement of how much time would be needed if all the threads' work is aggregated in a sequential manner instead of having any time overlap which threading actually provides.

Varying number of transactions

num_tx	exec_time (in secs)							
	naive_pthreads		smart_pthreads		naive_sequential_pthreads		smart_sequential_pthreads	
	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset
50	0.00234	0.001781	0.008364	0.007621	0.000138	0.000195	0.007085	0.007427
100	0.022403	0.014834	0.051744	0.040673	0.109489	0.119017	0.278568	0.196157
150	0.065483	0.062736	0.197969	0.157017	0.608572	0.59409	2.669247	1.48227
200	0.1907	0.18525	0.458289	0.456888	2.207017	2.118227	7.721343	6.737537
250	0.409843	0.401534	0.91913	0.942431	5.054997	4.971323	15.63273 3	14.5601
300	0.961019	0.946423	1.49761	1.551983	10.25983 3	10.1094	23.7267	23.50166 7
350	2.942077	2.89004	3.018837	3.259337	25.7541	24.681133	44.42923 3	46.91333 3
400	4.405237	4.367143	4.434817	5.05253	44.96326 7	44.12286 7	63.88483 3	71.8936
450	8.07115	8.02166	6.896717	9.15751	68.39916 7	67.3581	88.59966 7	114.59566 7
500	9.539257	9.493583	7.94909	11.363633	85.0115	84.05393 3	105.5813 33	144.5283 33

Varying number of items

num_items	exec_time (in secs)			
	naive_pthreads	smart_pthreads	naive_sequential_pthreads	smart_sequential_pthreads

	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset
10	0.007502	0.00119	0.014263	0.001721	0.000583	0.000641	0.001773	0.001621
20	0.012103	0.002637	0.020054	0.006084	0.006174	0.007258	0.01457	0.015081
30	0.023424	0.004747	0.034472	0.014736	0.021077	0.022852	0.051466	0.050786
40	0.019995	0.012194	0.042138	0.02959	0.056462	0.06043	0.151517	0.10602
50	0.02619	0.012056	0.050613	0.035323	0.09157	0.100593	0.223726	0.169966
60	0.027828	0.022885	0.069786	0.062365	0.199794	0.212304	0.412951	0.356435
70	0.041963	0.023368	0.085939	0.070525	0.297285	0.282892	0.673827	0.483343
80	0.045959	0.036711	0.145888	0.109074	0.529955	0.50452	1.86832	0.984555
90	0.049237	0.042574	0.181462	0.127912	0.64428	0.612966	2.354497	1.156083
100	0.065052	0.060287	0.329298	0.185341	1.128857	1.077317	5.986547	2.222963

Varying threshold

threshold	exec_time (in secs)							
	naive_pthreads		smart_pthreads		naive_sequential_pthreads		smart_sequential_pthreads	
	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset	affinity_ set	affinity_ unset
25	5.529357	5.479483	5.189437	5.85543	51.7116	50.77993 3	75.60696 7	84.7291
50	0.501456	0.496973	1.21638	1.192073	5.68478	5.633617	19.6284	17.111933
75	0.071761	0.071238	0.225797	0.207486	0.620428	0.610308	1.561743	1.089
100	0.047836	0.046896	0.180064	0.173473	0.471179	0.464163	0.858971	0.856926
125	0.036768	0.033892	0.149795	0.142819	0.165763	0.164941	0.417571	0.388114
150	0.01587	0.01533	0.102922	0.101526	0.000767	0.000813	0.102478	0.101515
175	0.014782	0.014736	0.100035	0.100693	0.000316	0.000323	0.100035	0.100693
200	0.014792	0.01474	0.105695	0.101807	0.000316	0.000321	0.105695	0.101807


Varying probability

probability	exec_time (in secs)							
	naive_pthreads		smart_pthreads		naive_sequential_pthreads		smart_sequential_pthreads	
	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset	affinity_set	affinity_unset
20%	0.000466	0.000471	0.000992	0.001033	0.000064	0.000055	0.000992	0.001033
30%	0.00214	0.002103	0.006143	0.006246	0.000065	0.000063	0.006143	0.006246
40%	0.018938	0.005166	0.031933	0.015904	0.012099	0.00802	0.026642	0.02414
50%	0.030355	0.013207	0.051678	0.0371	0.096459	0.107021	0.177514	0.176375
60%	0.144184	0.134687	0.441738	0.395971	1.094197	1.017757	6.563383	4.75761
70%	3.834503	3.76301	8.667087	12.372633	30.447267	30.1052	99.315233	144.856

Special datasets (skewness of item popularity)-

For analysis of all the datasets mentioned in Section 5.6., the threshold has been set to 20 (same as the baseline config), whereas the level has been set to 3 since the baseline results only generate up to a max level of 3.

Data skewness



The main observation here is that for all the variant datasets, the execution times are considerably higher than that of the baseline dataset. Along with this, there is a significant difference even in the execution times between different variant datasets. This indicates that the skewness pattern in the data also affects the required computation for the same workload size i.e. the distribution of data does matter. The smart implementation performing worse than the naive one is due to the reason mentioned previously (see the cpu affinity set and unset results). Though, again as expected, pthreads implementation performs better than the serial one.

Another interesting observation is that GPU performs much better compared to other implementations in all the skewness variants. This could be because all these variants support data locality in memory accesses as evident from the closeness of 1's as seen in the figures in Section 5.6.

The count of 1s in each of the compared datasets is as shown below-

Dataset	Count (out of total)
Baseline	1238 / 2500
Variant 1	1275 / 2500
Variant 2	1250 / 2500
Variant 3	1486 / 2500
Variant 4	1024 / 2500

For each of these datasets, the number of frequent itemsets generated is as follows-

Dataset	Total frequent itemsets
Baseline	62
Variant 1	4991
Variant 2	4525
Variant 3	10700
Variant 4	1350

Varying the number of cores-

The taskset command line utility has been used to limit the number of cores the process can use to run with. Accordingly, during each run, changes were made in the code to only allow those many as available cores.

```
taskset --cpu-list <cpu_range> ./driver
```


Example: `taskset --cpu-list 0-23 ./driver`

The parameters set were-

NUM_TX = 500
NUM_ITEMS = 50
PROBABILITY = 0.5
THRESHOLD = 20

Number of cores and which ones specifically were used is as shown below-

Number of cores	Core range
48	0 - 47
24	0 - 23
12	0 - 11
6	0 - 5
3	0 - 2



As seen in the graph above, the sequential implementation does not get affected by reducing the number of allowed cores since it anyways runs on just a single core. On the other hand, in the case of the pthreads implementation, when the number of allowed cores is decreased, there is a linear increase in the execution time. This behaviour is as expected since the number of cores available for compute are less and hence it'll take more time to execute the same workload. The smart implementation shows lower execution times than the naive implementation in the obtained results, which is as expected. Finally, as can be observed, the pthreads implementation performs better than the serial one.

Restricted CPU usage based upon sar output as set by tasklet-

```
7958 04:41:12 PM CPU      user      %nice    %system    %iowait    %steal    %idle
7959 04:41:14 PM all      20.55      0.00      6.49      0.00      0.00      72.96
7960 04:41:14 PM 0        59.18      0.00      7.14      0.00      0.00      33.67
7961 04:41:14 PM 1        47.47      0.00      2.02      0.00      0.00      50.51
7962 04:41:14 PM 2        37.11      0.00      10.31      0.00      0.00      52.58
7963 04:41:14 PM 3        29.17      0.00      13.54      0.00      0.00      57.29
7964 04:41:14 PM 4        22.00      0.00      5.00      0.00      0.00      73.00
7965 04:41:14 PM 5        39.00      0.00      9.00      0.00      0.00      52.00
7966 04:41:14 PM 6        31.25      0.00      12.50      0.00      0.00      56.25
7967 04:41:14 PM 7        41.24      0.00      10.31      0.00      0.00      48.45
7968 04:41:14 PM 8        44.33      0.00      7.22      0.00      0.00      48.45
7969 04:41:14 PM 9        57.00      0.00      5.00      0.00      0.00      38.00
7970 04:41:14 PM 10       34.69      0.00      10.20      0.00      0.00      55.10
7971 04:41:14 PM 11       39.18      0.00      9.28      0.00      0.00      51.55
7972 04:41:14 PM 12       17.35      0.00      9.18      0.00      0.00      73.47
7973 04:41:14 PM 13       52.53      0.00      5.05      0.00      0.00      42.42
7974 04:41:14 PM 14       53.06      0.00      7.14      0.00      0.00      39.80
7975 04:41:14 PM 15       24.24      0.00      4.04      0.00      0.00      71.72
7976 04:41:14 PM 16       34.38      0.00      8.33      0.00      0.00      57.29
7977 04:41:14 PM 17       40.48      0.00      6.06      0.00      0.00      45.45
7978 04:41:14 PM 18       40.48      0.00      7.22      0.00      0.00      43.30
7979 04:41:14 PM 19       38.78      0.00      10.20      0.00      0.00      51.02
7980 04:41:14 PM 20       15.31      0.00      7.14      0.00      0.00      77.55
7981 04:41:14 PM 21       29.00      0.00      1.00      0.00      0.00      70.00
7982 04:41:14 PM 22       37.09      0.00      9.47      0.00      0.00      52.63
7983 04:41:14 PM 23       36.08      0.00      11.34      0.00      0.00      52.58
7984 04:41:14 PM 24       0.00      0.00      0.00      0.00      0.00      100.00
7985 04:41:14 PM 25       37.37      0.00      62.63      0.00      0.00      0.00
7986 04:41:14 PM 26       0.00      0.00      0.00      0.00      0.00      100.00
7987 04:41:14 PM 27       0.00      0.00      0.00      0.00      0.00      100.00
7988 04:41:14 PM 28       0.00      0.00      0.00      0.00      0.00      100.00
7989 04:41:14 PM 29       0.00      0.00      0.00      0.00      0.00      100.00
7990 04:41:14 PM 30       0.00      0.00      0.00      0.00      0.00      100.00
7991 04:41:14 PM 31       0.00      0.00      0.00      0.00      0.00      100.00
7992 04:41:14 PM 32       0.00      0.00      0.00      0.00      0.00      100.00
7993 04:41:14 PM 33       0.00      0.00      0.00      0.00      0.00      100.00
7994 04:41:14 PM 34       0.00      0.00      0.00      0.00      0.00      100.00
7995 04:41:14 PM 35       0.00      0.00      0.00      0.00      0.00      100.00
7996 04:41:14 PM 36       0.00      0.00      0.00      0.00      0.00      100.00
7997 04:41:14 PM 37       0.00      0.00      0.00      0.00      0.00      100.00
7998 04:41:14 PM 38       0.00      0.00      0.00      0.00      0.00      100.00
7999 04:41:14 PM 39       0.00      0.00      0.00      0.00      0.00      100.00
8000 04:41:14 PM 40       39.39      0.00      60.61      0.00      0.00      0.00
8001 04:41:14 PM 41       0.00      0.00      0.00      0.00      0.00      100.00
8002 04:41:14 PM 42       0.00      0.00      0.00      0.00      0.00      100.00
8003 04:41:14 PM 43       0.00      0.00      0.00      0.00      0.00      100.00
8004 04:41:14 PM 44       0.00      0.00      0.00      0.00      0.00      100.00
8005 04:41:14 PM 45       0.00      0.00      0.00      0.00      0.00      100.00
8006 04:41:14 PM 46       0.00      0.00      0.00      0.00      0.00      100.00
8007 04:41:14 PM 47       0.00      0.00      0.00      0.00      0.00      100.00
8008 04:41:14 PM 48       0.00      0.00      0.00      0.00      0.00      100.00
8009
```



```
File Edit View Bookmarks Settings Help
affinity_results: vi --- Konsole
cpu_usage_affinity_unset_3_cores.txt
8723
8724 05:12:09 PM CPU %user %nice %system %iowait %steal %idle
8725 05:12:10 PM all 7.05 0.00 2.95 0.02 0.00 89.97
8726 05:12:10 PM 0 86.87 0.00 0.00 0.00 0.00 3.05
8727 05:12:10 PM 1 89.11 0.00 4.95 0.00 0.00 5.94
8728 05:12:10 PM 2 88.61 0.00 9.18 0.00 0.00 10.20
8729 05:12:10 PM 3 0.00 0.00 0.00 0.00 0.00 100.00
8730 05:12:10 PM 4 0.00 0.00 0.00 0.00 0.00 100.00
8731 05:12:10 PM 5 0.99 0.00 0.00 0.00 0.00 99.01
8732 05:12:10 PM 6 0.00 0.00 0.00 0.00 0.00 100.00
8733 05:12:10 PM 7 0.00 0.00 0.00 0.00 0.00 100.00
8734 05:12:10 PM 8 0.00 0.00 0.98 0.98 0.00 98.04
8735 05:12:10 PM 9 1.00 0.00 0.00 0.00 0.00 99.00
8736 05:12:10 PM 10 0.00 0.00 0.00 0.00 0.00 100.00
8737 05:12:10 PM 11 0.00 0.00 0.00 0.00 0.00 100.00
8738 05:12:10 PM 12 0.00 0.00 0.00 0.00 0.00 100.00
8739 05:12:10 PM 13 0.99 0.00 0.00 0.00 0.00 99.01
8740 05:12:10 PM 14 2.00 0.00 0.00 0.00 0.00 98.00
8741 05:12:10 PM 15 0.00 0.00 0.00 0.00 0.00 100.00
8742 05:12:10 PM 16 0.00 0.00 0.00 0.00 0.00 100.00
8743 05:12:10 PM 17 0.00 0.00 0.00 0.00 0.00 100.00
8744 05:12:10 PM 18 0.00 0.00 0.00 0.00 0.00 100.00
8745 05:12:10 PM 19 0.00 0.00 0.00 0.00 0.00 100.00
8746 05:12:10 PM 20 0.00 0.00 0.00 0.00 0.00 100.00
8747 05:12:10 PM 21 0.00 0.00 0.00 0.00 0.00 100.00
8748 05:12:10 PM 22 0.00 0.00 0.00 0.00 0.00 100.00
8749 05:12:10 PM 23 0.00 0.00 0.00 0.00 0.00 100.00
8750 05:12:10 PM 24 0.00 0.00 0.00 0.00 0.00 100.00
8751 05:12:10 PM 25 40.00 0.00 60.00 0.00 0.00 0.00
8752 05:12:10 PM 26 0.00 0.00 0.00 0.00 0.00 100.00
8753 05:12:10 PM 27 0.00 0.00 0.00 0.00 0.00 100.00
8754 05:12:10 PM 28 0.00 0.00 0.00 0.00 0.00 100.00
8755 05:12:10 PM 29 0.00 0.00 0.00 0.00 0.00 100.00
8756 05:12:10 PM 30 0.00 0.00 0.00 0.00 0.00 100.00
8757 05:12:10 PM 31 0.00 0.00 0.00 0.00 0.00 100.00
8758 05:12:10 PM 32 0.00 0.00 0.00 0.00 0.00 100.00
8759 05:12:10 PM 33 0.00 0.00 0.00 0.00 0.00 100.00
8760 05:12:10 PM 34 0.00 0.00 0.00 0.00 0.00 100.00
8761 05:12:10 PM 35 0.00 0.00 0.00 0.00 0.00 100.00
8762 05:12:10 PM 36 0.00 0.00 0.00 0.00 0.00 100.00
8763 05:12:10 PM 37 0.00 0.00 0.00 0.00 0.00 100.00
8764 05:12:10 PM 38 0.00 0.00 0.00 0.00 0.00 100.00
8765 05:12:10 PM 39 0.00 0.00 0.00 0.00 0.00 100.00
8766 05:12:10 PM 40 40.00 0.00 60.00 0.00 0.00 0.00
8767 05:12:10 PM 41 0.00 0.00 0.00 0.00 0.00 100.00
8768 05:12:10 PM 42 0.00 0.00 0.00 0.00 0.00 100.00
8769 05:12:10 PM 43 0.00 0.00 0.00 0.00 0.00 100.00
8770 05:12:10 PM 44 0.00 0.00 0.00 0.00 0.00 100.00
8771 05:12:10 PM 45 0.00 0.00 0.00 0.00 0.00 100.00
8772 05:12:10 PM 46 0.00 0.00 0.00 0.00 0.00 100.00
8773 05:12:10 PM 47 0.00 0.00 0.00 0.00 0.00 100.00
8774
NORMAL | 1 | 115,157 words | 67% | 8774/13059 | 1
```

As expected, in case of the sequential implementation, one of the cores among the restricted set is utilized (**note**: image not shown above).

5.8. Summarizing the linux utilities used

1. **sar**- can be used to monitor cpu usage characteristics (one advantage of this command is that its output can be redirected).
2. **htop** / **top**- similar to sar, but these commands provide a lot of other information like running processes, which users are running which processes, etc.
3. **taskset**- can be used to run any other command by limiting / restricting the allowed cores for that command.
4. **nvprof**- a profiling tool that enables collection and viewing of profiling data from the command line (a part of the CUDA toolkit).

6. Programming challenges

One of the biggest challenges for us was to figure out what data structure to use in the CUDA kernel. Since we used vectors in cpu implementations, we were first inclined to have a vector version in CUDA. However, we decided to go for a simple linear array as adopted in the general play-code/tutorial examples of vector addition.

Another challenge was to directly adopt the textbook's parallel algorithm in CUDA. We were not able to figure out how to do the computations of level 3 and later with just one cuda kernel launch, particularly, how would the GPU share intermediate results with the CPU for bookkeeping when required, and storing all the levels in a common data structure shared by all GPU threads. So, we used another approach as described in Section 3.5.

7. Hardware bottlenecks

1. All the cores in CPU are generally not fully available as we found some other processes running for long duration (like days).
2. Our GPU implementation could benefit from concurrent kernel launches. However, it is difficult to program, and usually kernels are executed sequentially by the GPU.
3. Unlike pthreads where a CPU thread can launch new threads dynamically to scale up when the computation for a thread increases, it is tricky for a GPU thread to efficiently launch another kernel.
4. The multicore analysis as shown in Section 5.7 suggests that the pthread implementation doesn't get benefit when going from 24 cores to 48 cores.

8. Limitations and future directions

1. More intelligent use of cpu affinity needs to be pursued. Currently, only a simple round robin runtime allocation mechanism has been used, which might not always be optimal since the results obtained do not show significant improvement with cpu affinity set, which ideally should be the case.
2. In the current pthreads implementation, there is parallelism introduced only at level 2's equivalence classes. But within each thread that handles this equivalence class, at further levels, all the computation happens in this single thread sequentially, though asynchronously. A further scope for optimization via more parallelism opportunities would be to launch new child threads for equivalence classes generated at each level, leading to a tree of threads.
3. An interesting way to find the least frequent items could be by inverting the matrix (meaning: setting all 1 to 0 and all 0 to 1). Sales experts could use this analysis to find certain items to remove from the store which could help in making space for new or popular items.
4. With a larger database size, and dense matrix, the GPU should show better speedups compared to a multi-core CPU. This can be determined as part of future study.
5. Writing efficient CUDA code - there are a lot of functions which CUDA programming supports to efficiently use GPUs. Our implementation was just a beginning endeavor to demonstrate the benefit of GPUs. However, we carefully understand this, and we aim to write efficient code which industry can adopt as part of their products.
6. Since the application heavily relies on reading and writing large arrays, there is a lot of scope on doing memory optimizations. Particularly, one can use the benefits of data locality, prefetching, and reducing memory divergence. Learning to profile parts of the program to identify the bottlenecks (API calls or long execution time code blocks) and correctly tuning the application should be a wise way to improve our application performance. We plan to follow this as part of future work.
7. Heterogenous style computing could be adopted based on the work to be done. If the work has less parallelization opportunity, it could be a good idea to run it on a CPU (which generally has limited cores). On the other hand, if the work has more

parallelization opportunities, it could be beneficial to prioritize it to run on the GPU. For eg: if the equivalence class has less than 20 items sets, run it on CPU, else run it on the GPU.

9. References

[1] ExampleApps CSE 532 Course Lecture Slides PPT

[2] Parallel Computer Architecture: A Hardware/Software Approach, Pgs 80 - 81, 179 - 181